



TECH TALK:

Architecture and Automation for Researchers

Friday, October 23, 2020
10:30 AM US EDT

Rajib Ghosh
Global Senior Solutions Architect
Oracle for Research

TECH TALK HOUSEKEEPING

- Today's webinar is being recorded. We will share the link to the recording with you via email after the event. The recording will also be made available to the Oracle for Research community.
- We invite your comments and questions, both about the tech topic being discussed and about the series more generally. Questions may be submitted using the Q&A box on your screen or you may ask questions directly using your microphone. When not asking a question, please mute your microphone.
- Questions may be asked during the presentation and we will also have a Q & A time at the end of the presentation when you can ask questions directly and engage in discussion.
- At Oracle for Research, we believe that research and innovation happen best when a diverse and thoughtful community is free to engage in respectful, compassionate, and open dialog. To that end, when asking a question or providing feedback, we ask that all participants be respectful, collaborative, and constructive.

Agenda

Recap and asks

1. Shape selection , storage and performance testing
2. Architectures and automation
3. Performance benchmarks & Data science (platform vs Image usage)

Architectures for Researchers

1. Standard researcher architecture (Demo)
2. Cloud bursting architectures

OCI Automation

1. Instance scaling (demo)
2. Resource stacks and Terraform (demo)
3. Programmatic automation with OCI CLI and API


Github and links

1. Technical How-Tos
2. Repositories and Images
3. Researcher collaboration

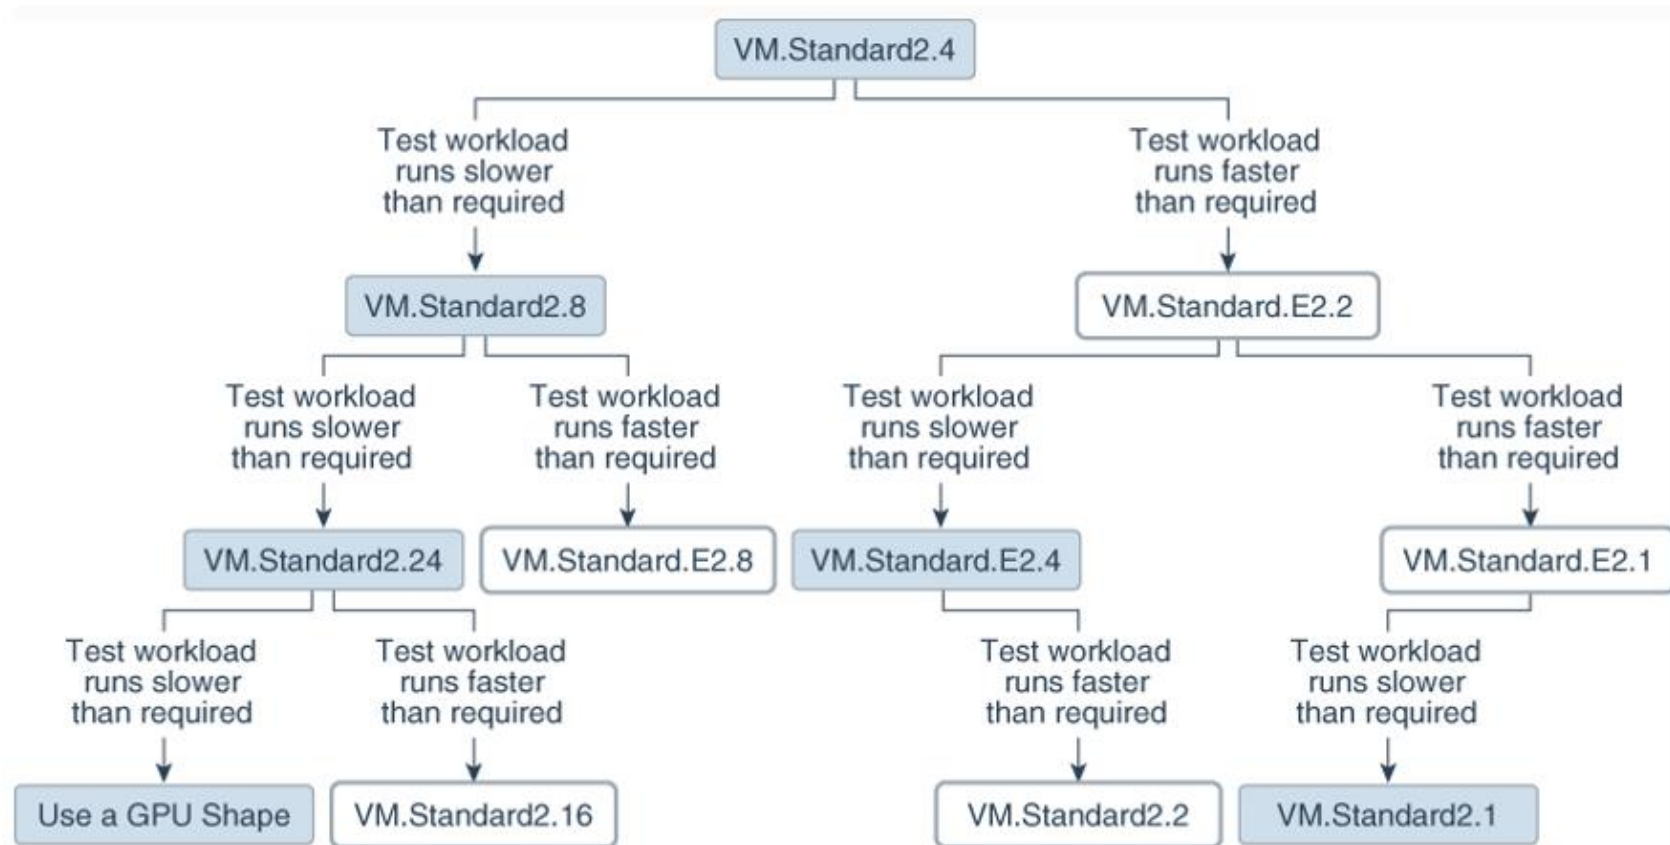
Q & A

Q & A
What works best? Any researcher wish list?

A GENERIC WORKFLOW TO DECIDE OPTIMAL SHAPES




SHAPE TESTING HIERARCHY FOR NON-GPU USAGE USE-CASE



NOTE : This should be used as a guideline and may differ based on researcher data and computation scenario

Courtesy : Oracle architecture center

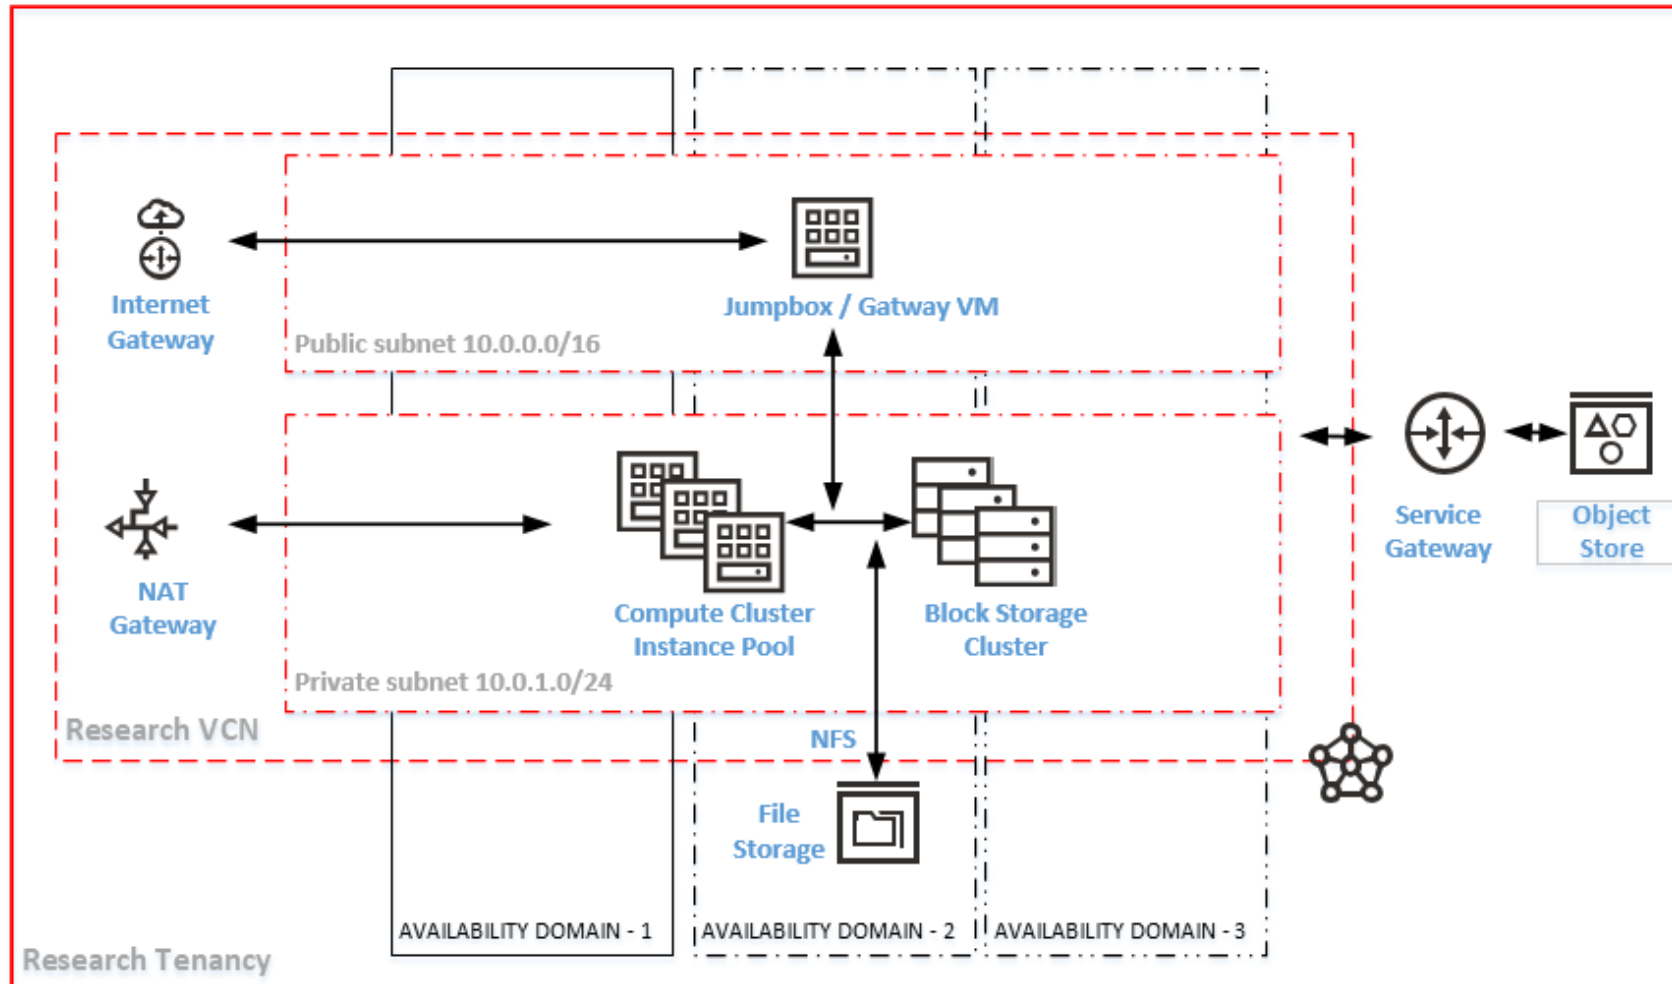
SHAPE TESTING HIERARCHY FOR GPU USAGE USE CASE



NOTE : This should be used as a guideline and may differ based on researcher data and computation scenario

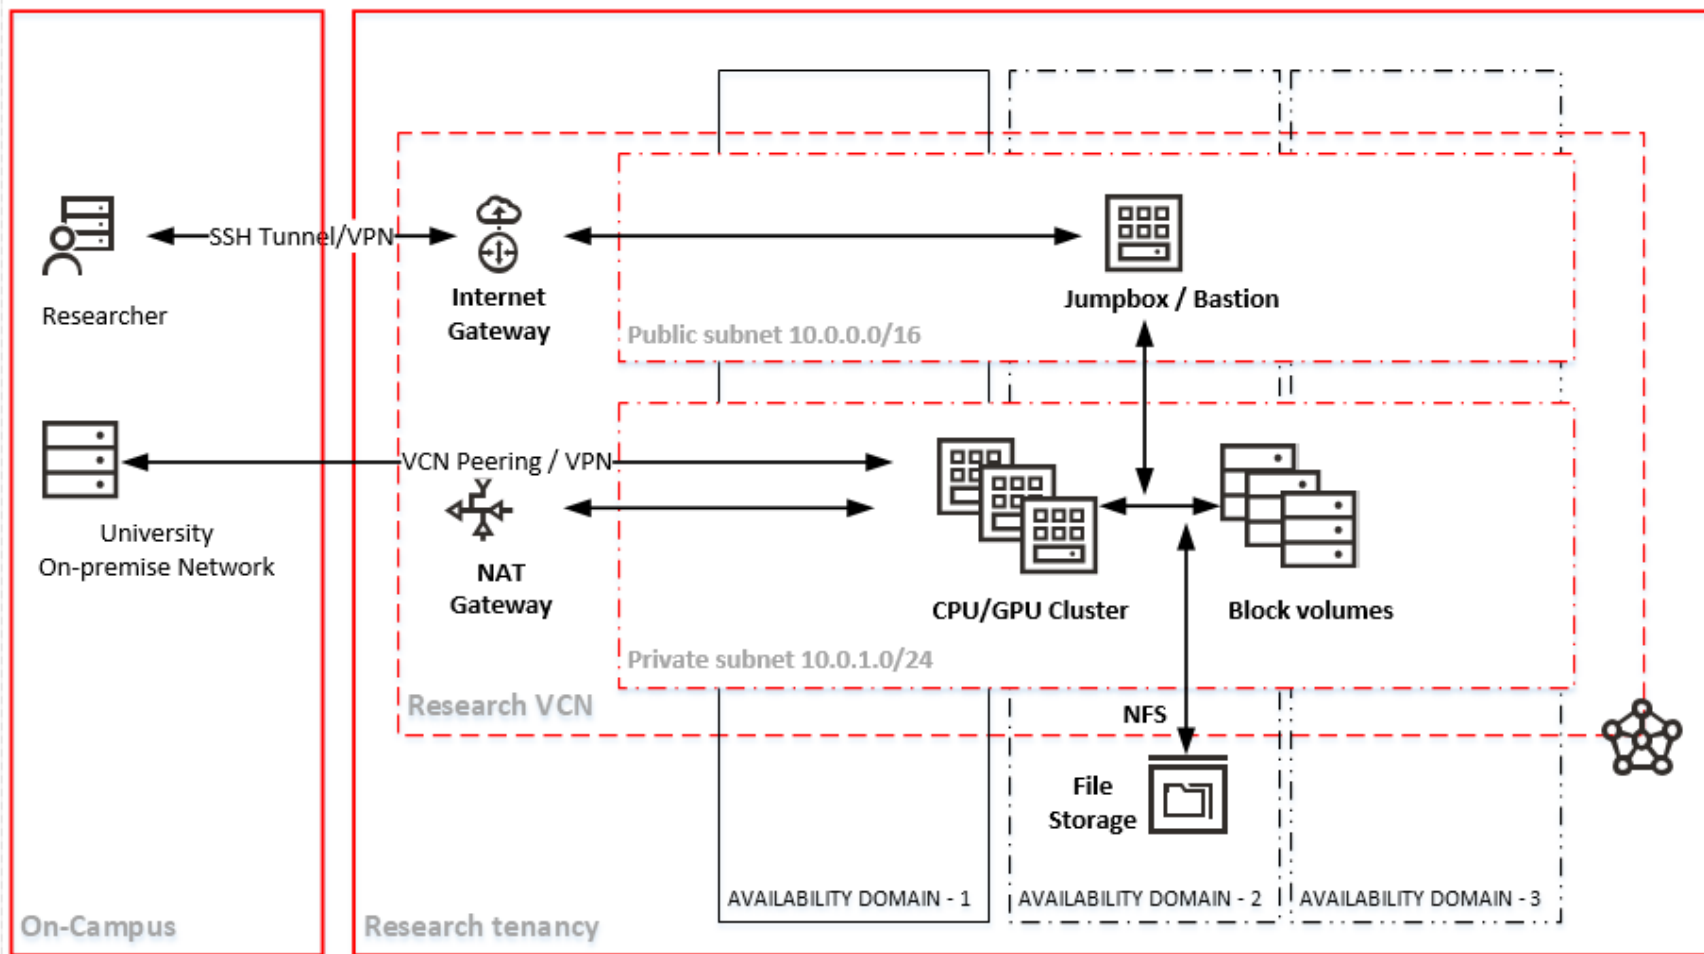
Courtesy : Oracle architecture center

OCI Standard cluster architecture for Researchers



Component	Recommendation
Jumpbox VM	Use Free-Tier VM
CPU/GPU Cluster	Manual Build + Block storage
CPU/GPU Cluster	Automated Build + Block storage
HPC Cluster	Instance Pool + RDMA
NAT Gateway	Download software to pvt subnet
File storage	For cross-AD exports

Cloud bursting architecture for Researchers



Key features

1. OCI access – (SSH Tunnel to Bastion to GPU or VPN)
2. CPU/GPU bursting (OCI Scheduling) - OCI Instance pooling + auto-scaling
3. CPU/GPU bursting (Campus Scheduler) - OCI CLI + Stacks or OCI API
4. NAT Gateway - Download software from public repositories
5. File storage - Cost effective data sharing within or across AD

Architecture use-cases

Standard Architecture

Recommended when

1. Quickly standup the application / cluster in Oracle cloud
2. Entire application would run in cloud
3. Database centric ML applications (Autonomous DB / Roracle)
4. Dynamic scaling of cluster nodes
5. Hybrid cluster with multiple shapes
6. Different workload types (Emb parallel / Tightly-coupled)
7. Simple SSH but secure access to computational VMs
8. All data contained in the cloud
9. Benchmarking against on-campus / clouds
10. Using Oracle cloud automation framework & scripts

Not recommended when

1. Computation in cloud but data is resident outside Oracle cloud
2. On-campus resources are sufficient for workloads
3. Non-research operational projects
4. Frequent transfer of code and data in/out from Oracle cloud

Cloud bursting Architecture

Recommended when

1. A percentage of the workload can be run on Oracle cloud
2. Infrequent cloud usage
3. On-premise workload scheduling
4. On-premise cluster is overloaded
5. Reduce workload by moving an application/data to OCI
6. Smaller computational data sets
7. Performance benchmarking against on-campus / other clouds
8. Scheduling automation is tested and working
9. Containerized applications & canary testing
10. Benchmarking against OCI shapes
11. Embarrassingly parallel HPC jobs

Not recommended when

1. Large data migration during bursting
2. Tightly-coupled apps (Cross node communication)

Why standard architecture is important?

Being Standard

1. Quick and easy to implement
2. Easier resource management
3. Single & common reference for all researchers

Automation

1. One-click standup and shutdown of instance through console (Stacks and Terraform)
2. Quick and easy configuration management
3. Programmatic automation (CLI / API) through a Free-tier VM

Better resource utilization

1. Planned utilization of CPU/GPU/Storage resources
2. Effective cluster and computational planning
3. Better credit usage and service limit allocation

Better benchmarking

1. Benchmarking against similar hardware & computational specifications
2. Leverage performance tips from Oracle
3. Testing same workload against higher end shapes

Quicker support

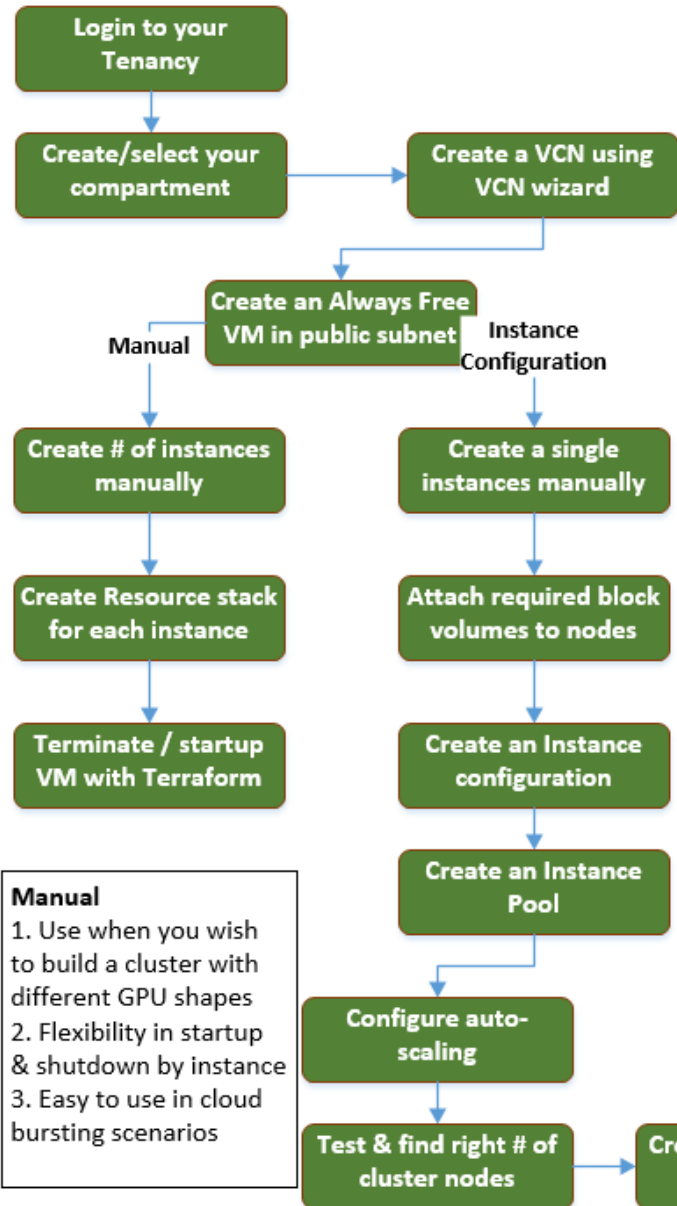
1. Issues can be resolved faster
2. Issues can be replicated/tested without access to researcher tenancies
3. Comparative benchmarks from the community can help tune your implementation

Security

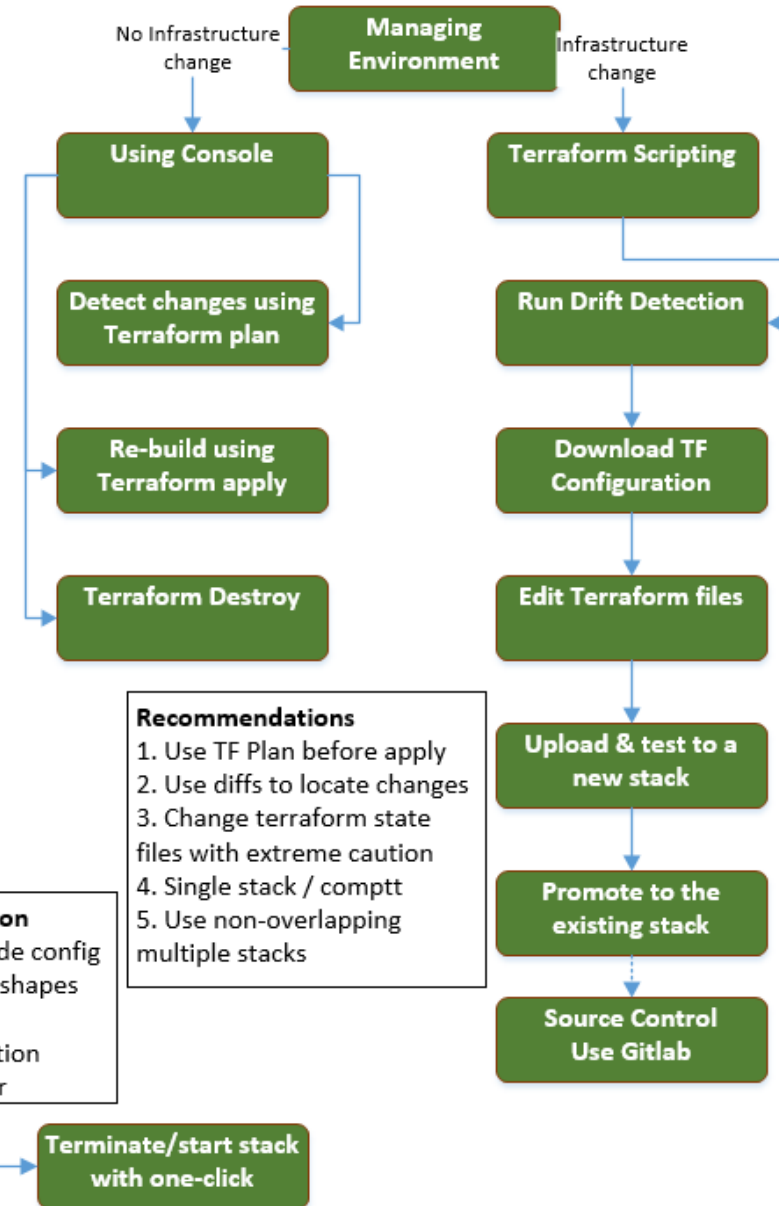
1. Private subnet insulates your computation VM and data
2. Access only through a jump box / secure gateway VM
3. NAT access provides direct internet access for code / patch downloads

Automation Setup & management workflow

Initial Creation



Managing Changes



Oracle for Research github – <https://github.com/OracleForResearch>

Technology How-Tos	<ol style="list-style-type: none">1. Short 2-page How-Tos (Categorized)2. Important OCI links (Supplemental pages)3. Aligned and dedicated to researcher needs	
Image Sandbox	<ol style="list-style-type: none">1. AIML Sandbox & AFNI (Base reference images with CPU/GPU versions on OL7 and Ubuntu)2. Actively worked on by Oracle for Research3. OFR will develop more based on active researchers participation and feedback	
Researcher Collaboration	<ol style="list-style-type: none">1. Contributions (benchmarks, test results, data)2. OFR and OCI reviews3. Researcher publications repository	
Images & Applications	<ol style="list-style-type: none">1. One stop place for Oracle provided images2. Oracle cloud Images for Research only3. Researcher image contribution	<ol style="list-style-type: none">1. Molecular dynamics2. Open source and proprietary images3. Agro & farm data images
Technology Talks	<ol style="list-style-type: none">1. Oracle for Research presentations2. Relevant Tech presentations for researchers3. Product updates	
Coming up..	<p>Benchmarking and guidelines Choosing the right data science platform. Researcher images for sharing and publication What we need? – Your feedback / what makes sense?</p>	



TECH TALK:

Architecture and Automation for Researchers

Questions, Answers & Discussion



TECH TALK:

Architecture and Automation for Researchers

Questions? Comments? Feedback?

Contact us!

Website: oracle.com/oracle-for-research/


Twitter: @OracleResearch

Email: OracleForResearchTech_ww@oracle.com

Next Tech Talk: November 13, 2020, 10:30AM EDT

Parking lot

For Reference



Oracle for Research Tech Talk Oracle Cloud Topics

Foundational	Oracle cloud – Getting you started and running Cloud instances and cloud storage options Migrating data and running computations
Architecture	Reference architecture patterns for researchers New features updates and recommended practices Performance benchmarks and data
Tools and automation	Tool selection, version and guidance Image repositories, Terraform and interfaces
HPC and cluster	High performance computing, workload classification, parallelization Cluster setup, utilization and monitoring
Machine and Deep Learning	Model selection guidance
Researcher guidance	Functional and data guidance and curation support Industry models with research computing

Instance Type	Shape series	Shape	Purpose
Virtual	Always Free	VM.StandardE2.1Micro	Automation control, gateway, configurations
	Standard	VM.Standard1.1~1.16	Low workload testing / Image builds / installs
	AMD (Gen 2)	VM.StandardE2.1~2.8	Prototype workload testing
	DenseIO	VM.DenseIO2.x (NVMe)	Heavy IO workload testing
	GPU (P100)	VM.GPU2.1	AI / ML or other GPU prototype testing
	GPU (V100)	VM.GPU3.1~3.4	Tensor core AI / DL workloads
	Intel Skylake (Fixed)	VM.Standard2.1~2.24	Workloads to save on credits
	AMD Rome (Flex)	VM.StandardE3.Flex	Benchmarking / price-performance
Bare metal	HPC	BM.HPC2.36 (NVMe)	CPU+high throughput for HPC workloads
	AMD (Gen 3)	BM.StandardE3.128	High CPU/throughput workloads
	Standard	BM.Standard1.36/B1.44	Low CPU/RAM utilization at lowest BM cost
	AMD (Gen 2)	BM.StandardE2.52	Best price-performance for BM workloads
	AMD (Gen 3)	BM.StandardE2.64	Best Gen3 price-performance for BM workloads
	DenseIO	BM.DenseIO2.52 (NVMe)	Best price performance for IO intensive workloads
	GPU (P100)	BM.GPU2.2	Benchmarking pascal based GPU workloads
	GPU (V100)	BM.GPU3.8	Best price performant for large GPU workloads
	GPU (A100)	BM.GPU4.8	Fastest GPU – large DL applications (pre-GA)

Oracle cloud storages

Storage type	Features	Recommended usage
On Campus storage	Good for on-campus data processing requirements Data could be distributed (laptops) or centralized Quick data retrieval for on-campus computations	Store data for on-campus computational purpose Store data if storage is available and it is already paid for
Object storage archive	Unlimited data storage in Oracle cloud Low cost (\$0.0026/GB/Month) by consumption Better if cost/GB is lower than on-campus storage	Not so good for frequent large data retrievals Not readily available for computational purpose
Object storage standard	Unlimited data storage with faster access than OS archive <u>Costs (0.0255/GB/Month) by consumption</u> Good for <u>frequently accessed data across cloud tenancies</u> Secured and encrypted data at rest and in transit (https)	Store data backups for quick downloads in cloud / campus Store large data volumes at relatively lower cost Store data that is infrequently processed
Block volume	Most common storage for compute/databases Mountable across multiple instances within an AD <u>Cost (0.0255/GB/month) – based on total volume in GB</u> Supports parallel filesystem & <u>best price performance</u>	Leverage to store computational data for most loads Extend storage / instance as needed Fill up allocated capacity to save on costs
File system storage	NFSv3 unlimited file system storage mountable across AD Higher cost (0.3/GB/month) <u>on consumption</u> Good for file sharing across tenancies and OS Performs linearly or better with higher data set size	Not good for cross AD data transfer in computational cycle Good for moving large volumes of data quickly across AD Use sparingly
Local NVMe	Highest IOPS and throughput & good for IO intensive loads Higher cost (built into compute) Non-persistent data Part of DenseIO and HPC shapes	Use for IO intensive parallel CPU workloads

Useful links

Getting Started

[Key concepts and terminology](#)

[Signing in to console, Sign-in options and changing your password](#)

[Setting up your tenancy](#)

[Tutorial – Launching your first Linux instance](#)

[Tutorial – Launching your first windows instance](#)

[Object storage and Pre-authenticated access](#)

[Image import and export](#)

[File storage system concepts](#)

[OCI Hands on labs](#)

[New features and navigation updates](#)

[Oracle cloud Free tier and FAQ](#)

[Custom key generation with puttygen or ssh-keygen](#)

[Frequently asked Questions](#)

[Getting help and contacting support](#)

Identity federation

[Federated identity for single sign-on](#)

[Migrating from on-premise to Oracle identity service](#)

Databases and moving data

[Oracle Autonomous databases and Tools](#)

[MySQL and NoSQL Services](#)

[Key concepts and terminology](#)

[Migrating databases to cloud](#)

[Loading data to autonomous with OCI Functions](#)

[Single-click move to autonomous](#)

Data science and AI/ML

[Oracle Data science platform and Tutorials](#)

[Genome analysis toolkit](#)

[Julia AI/HPC GPU Image](#)

[NVIDIA images and NVIDIA GPU image](#)

[Building a ML sandbox on Oracle cloud](#)

[Setting up an open-source ML and AI Environment](#)

[Machine learning autonomously](#)

High performance computing (HPC)

[Oracle HPC Cluster and Oracle HPC File system](#)

[NVIDIA GPU Cloud machine image](#)

[Oracle Linux 7 Cluster Networking Image](#)

[Oracle marketplace slurm image \(HPC + Slurm combo\)](#)

[Oracle cloud slurm image](#)

[Github OCI-HPC](#)

[Enabling HPC Cluster networking](#)

[Deploy High performance computing on Oracle cloud](#)

[Infrastructure](#)

[Deploy scalable and distributed file system using Lustre](#)

[Deploy BEEGFS parallel file system](#)

[UoB Cluster in the cloud](#)

[Cluster in the cloud - github](#)

[Molecular dynamics NAMD runbook](#) and [GROMACS runbooks](#)

Usage, billing and credit control

[Oracle cloud storage costs](#)

[Resource billing for stopped instances](#)

[Oracle cloud universal credit PaaS and IaaS service descriptions](#)

Oracle for Research Tech Talk Series Will Cover:

Technology training

- Reference architectures
- Best practices
- Tools and automation
- Cost control

Collaboration

- Q & A Live discussions
- Tips and tricks library
- Community forum
- Meet Oracle experts

Product announcements

- OFR technology updates
- OCI product updates
- Images and containers
- Public data

Researcher for researchers

- Technology Innovations
- Researcher publications
- Benchmarks
- Lessons learned

OCI Automation for Researchers

Common Researcher Issues	<ol style="list-style-type: none">1. Where to start?2. What shapes, images and storages do I need?3. How to manage my credit allocation effectively? <ol style="list-style-type: none">1. Any automation to terminate idle instances?2. How do I scale or burst my workload?3. Can I use my on-campus identity to login?
OCI Architectures for Researchers	<ol style="list-style-type: none">1. Standard researcher architecture – Bastion Free Tier VM + Private subnet for compute2. Shape and storage selection guidelines3. Scaling and cloud bursting architectures
Automation and usage control	Instance scaling (demo) Resource stacks and Terraform (demo) Tooling with OCI CLI
Credit control mechanisms	Cost analysis and cost reports Setting budgets and alerts
Github and links	Overview of the repositories Collaboration with Oracle and Researchers
Q & A	Q & A What works best and researcher wish list?